

# Computer Organization & Architecture

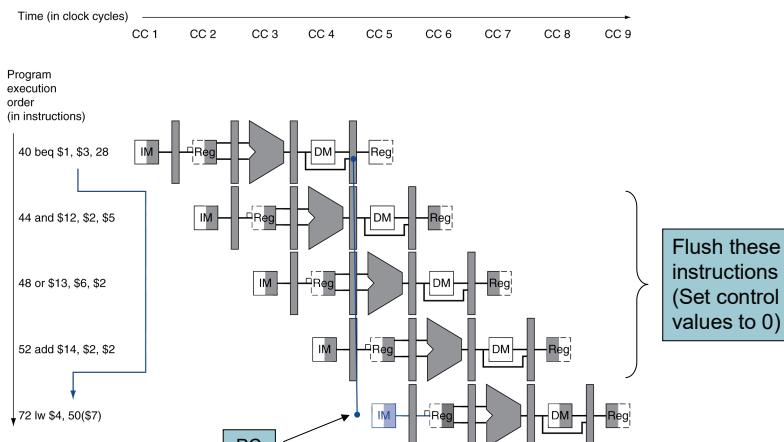
## 计算机组成与结构

### Chapter 4.8 Branch Prediction

Dr. Qingfeng (Karen) Zhuge, Professor  
 College of Computer Science and Software Engineering  
 East China Normal University

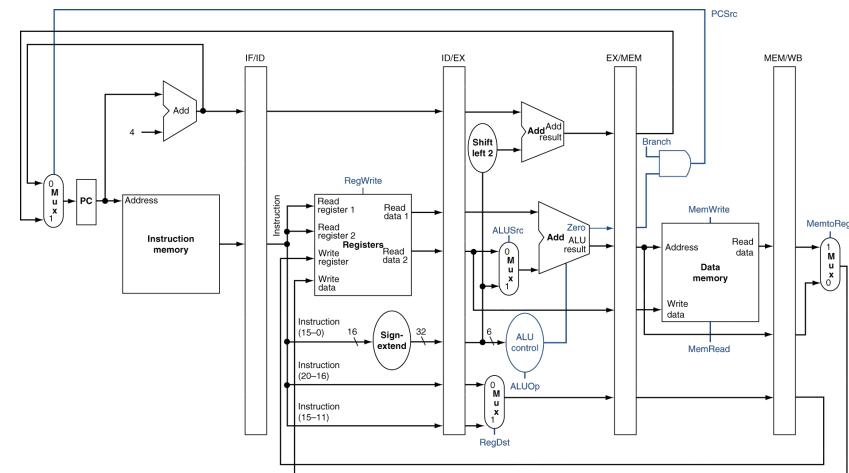
## Branch Hazards

- If branch outcome determined in MEM



Chapter 4 — The Processor — 3

## Pipelined Control (Simplified)



Chapter 4 — The Processor — 2

## Reducing Branch Delay

- Move hardware to determine outcome to ID stage

- Target address adder
- Register comparator

- Example: branch taken

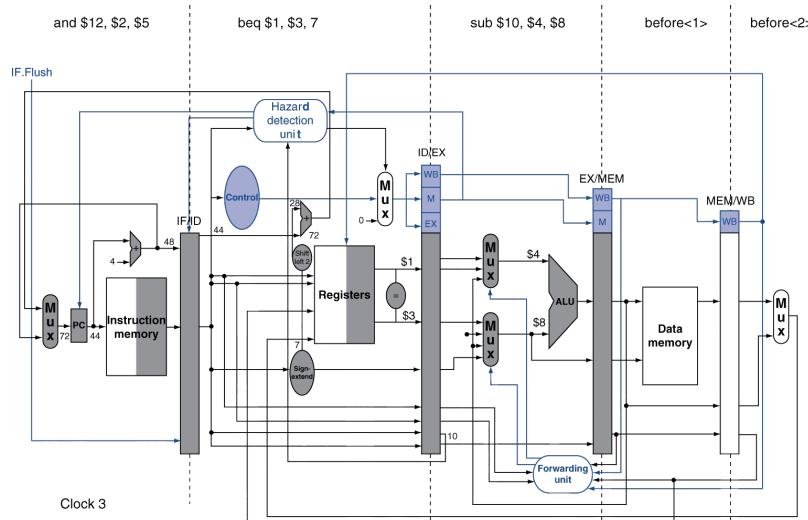
```

36: sub $10, $4, $8
40: beq $1, $3, 7
44: and $12, $2, $5
48: or $13, $2, $6
52: add $14, $4, $2
56: slt $15, $6, $7
...
72: lw $4, 50($7)

```

Chapter 4 — The Processor — 4

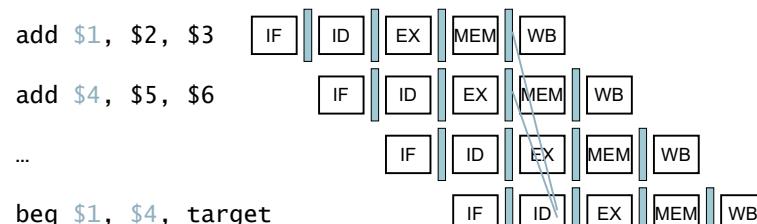
## Example: Branch Taken



Chapter 4 — The Processor — 5

## Data Hazards for Branches

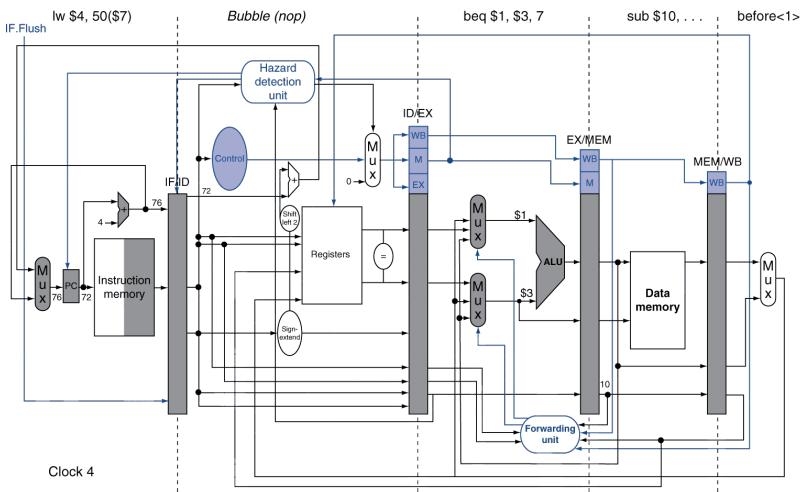
- If a comparison register is a destination of 2<sup>nd</sup> or 3<sup>rd</sup> preceding ALU instruction



- Can resolve using forwarding

Chapter 4 — The Processor — 7

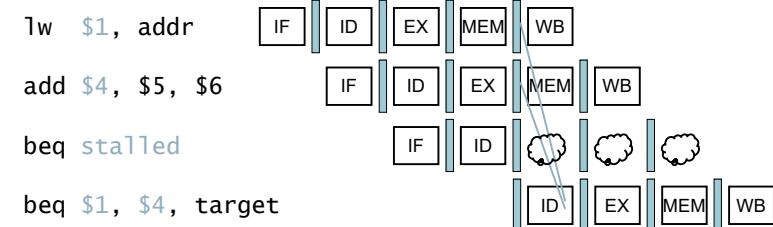
## Example: Branch Taken



Chapter 4 — The Processor — 6

## Data Hazards for Branches

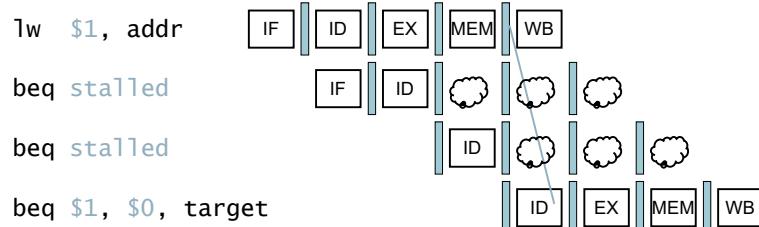
- If a comparison register is a destination of preceding ALU instruction or 2<sup>nd</sup> preceding load instruction
  - Need 1 stall cycle



Chapter 4 — The Processor — 8

## Data Hazards for Branches

- If a comparison register is a destination of immediately preceding load instruction
  - Need 2 stall cycles



Chapter 4 — The Processor — 9



## Reducing Pipeline Branch Penalties - I: Freezing (Flushing) the Pipeline

- Assuming now that the target address calculation and the comparison are performed during the ID stage, then, in case of "branch taken", the PC is changed at the end of the ID stage
  - simplest strategy:** freezing the pipeline as soon as the instruction is detected as a branch and regardless of its outcome (**branch penalty is fixed, SW cannot reduce it**)

| instruction                  | 1  | 2  | 3        | 4        | 5        | 6        | 7        | 8   | 9  |
|------------------------------|----|----|----------|----------|----------|----------|----------|-----|----|
| <code>BNEZ R5, target</code> | IF | ID |          |          |          |          |          |     |    |
| fall-through inst            |    | IF | (circle) | (circle) | (circle) | (circle) | (circle) |     |    |
| branch target                |    |    | IF       | ID       | EX       | MEM      | WB       |     |    |
| branch target + 1            |    |    |          | IF       | ID       | EX       | MEM      | WB  |    |
| branch target + 2            |    |    |          |          | IF       | ID       | EX       | MEM | WB |

## Branch (Control) Hazards and Their Impact on the Pipeline Performance

- Control hazards can cause a greater performance loss for a pipelined implementation than data hazards. In first approximation:

$$\text{speedupFromPipelining} = \frac{\text{pipelineDepth}}{1 + [\text{branchFrequency} \times \text{branchPenalty}]}$$

- performance loss can vary between 10% and 30% depending on the branch frequency
- generally, the deeper the pipeline, the worse the branch penalty (measured in clock cycles)

- Various strategies to reduce the branch penalty
  - static (compile time) schemes
    - they are fixed for each branch during the entire execution
  - dynamic
    - hardware and software techniques
    - more sophisticated branch prediction schemes

## Reducing Pipeline Branch Penalties - II: Predicting "Non-Taken"

- If the branch is taken, all the fetched instruction must be turned into NOPs and it is necessary to restart the fetch at the target address. The state must be unaffected!!

| instruction           | 1  | 2  | 3        | 4        | 5        | 6        | 7        | 8   | 9  |
|-----------------------|----|----|----------|----------|----------|----------|----------|-----|----|
| <b>untaken branch</b> | IF | ID |          |          |          |          |          |     |    |
| Fall-Through Inst.    |    | IF | ID       | EX       | MEM      | WB       |          |     |    |
| FTI + 1               |    |    | IF       | ID       | EX       | MEM      | WB       |     |    |
| FTI + 2               |    |    |          | IF       | ID       | EX       | MEM      | WB  |    |
| FTI + 3               |    |    |          |          | IF       | ID       | EX       | MEM | WB |
| instruction           | 1  | 2  | 3        | 4        | 5        | 6        | 7        | 8   | 9  |
| <b>taken branch</b>   | IF | ID |          |          |          |          |          |     |    |
| Fall-Through Inst.    |    | IF | (circle) | (circle) | (circle) | (circle) | (circle) |     |    |
| branch target         |    |    | IF       | ID       | EX       | MEM      | WB       |     |    |
| branch target + 1     |    |    |          | IF       | ID       | EX       | MEM      | WB  |    |
| branch target + 2     |    |    |          |          | IF       | ID       | EX       | MEM | WB |

## Delayed Branches

- ❑ If the branch hardware has been moved to the ID stage, then we can eliminate all branch stalls with **delayed branches** which are defined as always executing the next sequential instruction after the branch instruction – the branch takes effect **after** that next instruction
  - MIPS compiler moves an instruction to immediately after the branch that is not affected by the branch (a **safe** instruction) thereby **hiding** the branch delay
- ❑ With deeper pipelines, the branch delay grows requiring more than one delay slot
  - Delayed branches have lost popularity compared to more expensive but more flexible (dynamic) hardware branch prediction
  - Growth in available transistors has made hardware branch prediction relatively cheaper

CSE431 Chapter 4B.13

Irwin, PSU, 2008

## Basic Branch Prediction: Branch-History Table (or Branch-Prediction Buffer)

### Branch-History Table

- 1-bit prediction scheme
  - remembers last outcome of some branches
- implemented as a small memory indexed by a lower portion (some least-significative bits) of the address of the branch instruction
  - like a direct-map cache where an access is always a hit
- useful only to reduce the branch delay when longer than the time to compute the possible target PC
- We don't even know if prediction is correct. It is just a hint (if wrong the bit is inverted and we restart from the correct PC)
  - e.g., branch instructions at the following addresses share the BHT entry
    - 00F2BC 0101 1100
    - 010A5D 1001 1100

| Index | Taken? |
|-------|--------|
| 0000  | 1      |
| 0001  | 0      |
| 0010  | 1      |
| 0011  | 0      |
| 0100  | 1      |
| 0101  | 0      |
| 0110  | 1      |
| 0111  | 1      |
| 1000  | 1      |
| 1001  | 0      |
| 1010  | 1      |
| 1011  | 1      |
| 1100  | 0      |
| 1101  | 1      |
| 1110  | 1      |
| 1111  | 1      |

## Scheduling Branch Delay Slots

### A. From before branch

```
add $1,$2,$3
if $2=0 then
  delay slot
```

```
becomes
```

```
if $2=0 then
  add $1,$2,$3
```

### B. From branch target

```
sub $4,$5,$6 ←
add $1,$2,$3
if $1=0 then
  delay slot
```

```
becomes
```

```
add $1,$2,$3
if $1=0 then
  sub $4,$5,$6
```

### C. From fall through

```
add $1,$2,$3
if $1=0 then
  delay slot
```

```
becomes
```

```
add $1,$2,$3
if $1=0 then
  sub $4,$5,$6
```

❑ A is the best choice, fills delay slot and reduces IC

❑ In B and C, the **sub** instruction may need to be copied, increasing IC

❑ In B and C, must be okay to execute **sub** when branch fails

CSE431 Chapter 4B.14

Irwin, PSU, 2008

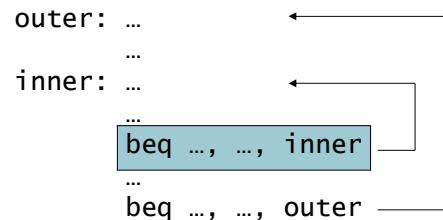
## Dynamic Branch Prediction

- In deeper and superscalar pipelines, branch penalty is more significant
- Use dynamic prediction
  - Branch prediction buffer (aka branch history table)
  - Indexed by recent branch instruction addresses
  - Stores outcome (taken/not taken)
  - To execute a branch
    - Check table, expect the same outcome
    - Start fetching from fall-through or target
    - If wrong, flush pipeline and flip prediction



## 1-Bit Predictor: Shortcoming

- Inner loop branches mispredicted twice!



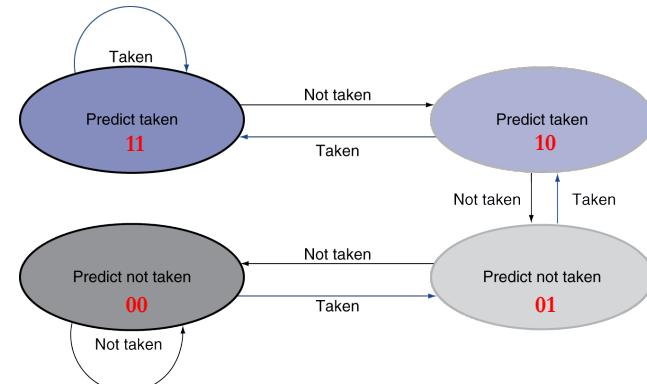
- Mispredict as taken on last iteration of inner loop
- Then mispredict as not taken on first iteration of inner loop next time around



Chapter 4 — The Processor — 17

## 2-Bit Predictor

- Only change prediction on two successive mispredictions



Chapter 4 — The Processor — 19

## Example: Accuracy of 1-Bit Prediction Scheme in the Presence of "Loop Branching"

**loop:**

```

L.D F0, 0(R1)
MUL.D F4, F0, F2
S.D F4, 0(R1)
DADDIU R1, R1, #-8
BNEZ R1 loop

```

- Assumptions
  - R1 is initialized to #80

| Iteration          | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
|--------------------|---|---|---|---|---|---|---|---|---|---|----|---|---|---|---|---|---|---|---|---|---|----|
| Predicted behavior | N | T | T | T | T | T | T | T | T | T | N  | T | T | T | T | T | T | T | T | T | T | T  |
| Actual behavior    | T | T | T | T | T | T | T | T | T | N | T  | T | T | T | T | T | T | T | T | T | N |    |

- Branch taken 90% of the times, but branch prediction accuracy is only 80%
- A 1-bit predictor for "loop branches" mispredicts at twice the rate that the branch is not taken

## Example: 2-Bit Prediction Scheme in Action with "Loop Branching"

**loop:**

```

L.D F0, 0(R1)
MUL.D F4, F0, F2
S.D F4, 0(R1)
DADDIU R1, R1, #-8
BNEZ R1 loop

```

- Assumptions
  - R1 is initialized to #80

| Iterations & steps | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
|--------------------|---|---|---|---|---|---|---|---|---|---|----|---|---|---|---|---|---|---|---|---|---|----|
| Predicted behavior | T | T | T | T | T | T | T | T | T | T | T  | T | T | T | T | T | T | T | T | T | T | T  |
| Actual behavior    | T | T | T | T | T | T | T | T | T | N | T  | T | T | T | T | T | T | T | T | T | T | N  |

- Branch taken 90% of the times, and branch prediction accuracy is now 90%
- The 2-bit predictor mispredicts at the 10<sup>th</sup> step of the 1<sup>st</sup> iteration, but *doesn't change its mind*. It just moves to "weak-predict-taken" for the 1<sup>st</sup> step of the following iteration

# Calculating the Branch Target

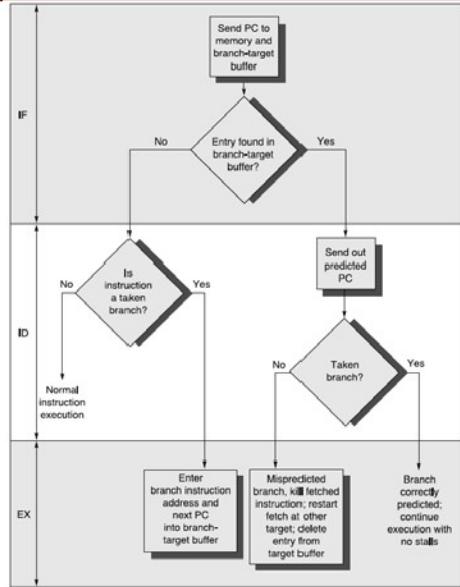
- Even with predictor, still need to calculate the target address
  - 1-cycle penalty for a taken branch
- Branch target buffer
  - Cache of target addresses
  - Indexed by PC when instruction fetched
    - If hit and instruction is branch predicted taken, can fetch target immediately



Chapter 4 — The Processor — 21

## Predicting the Instruction Target Address: Branch-Target Buffer (or Branch-Target Cache)

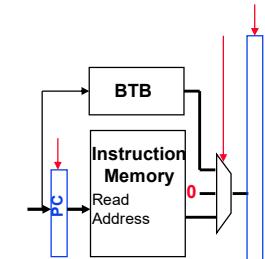
- Goal: learn the predicted instruction address at the end of IF stage
  - one cycle earlier
    - at best, branch-prediction buffers know the next predicted instruction at the end of ID
- No branch delay
  - if entry found in buffer and prediction is correct
  - if entry not found and branch is not taken
- Two cycle penalty
  - if prediction is incorrect
  - if entry is not found and branch is taken



## Branch Target Buffer

- The BHT predicts *when* a branch is taken, but does not tell *where* its taken to!
  - A **branch target buffer (BTB)** in the IF stage caches the branch target address, but we also need to fetch the next sequential instruction. The prediction bit in IF/ID selects which "next" instruction will be loaded into IF/ID at the next clock edge
    - Would need a two read port instruction memory
  - Or the BTB can cache the branch taken **instruction** while the instruction memory is fetching the next sequential instruction
- If the prediction is correct, stalls can be avoided no matter which direction they go

CSE431 Chapter 4B.22



Irwin, PSU, 2008

## Branch Target Buffer: Penalty Table

| Instruction in buffer | Prediction | Actual Branch | Penalty Clock Cycles |
|-----------------------|------------|---------------|----------------------|
| yes                   | taken      | taken         | 0                    |
| yes                   | taken      | not taken     | 2                    |
| no                    |            | taken         | 2                    |
| no                    |            | not taken     | 0                    |

- Assuming
  - 90% buffer hit rate, 85% prediction accuracy, 60% branches taken
- $P(\text{branch in buffer but not taken}) = 0.9 \times 0.15 = 0.135$
- $P(\text{branch not in buffer, but taken}) = 0.1 \times 0.6 = 0.06$
- $\text{Total branch penalty} = (0.135 + 0.06) \times 2 = 0.39$ 
  - the penalty for delayed-branch was about 0.5 clock cycles
  - penalty is lower with better predictors (and bigger branch delays)

## Adding "Global" Information to Branch Prediction: Correlating (or Two-Level) Branch Predictors

```

if (a == 2)
    a = 0;
...
if (b == 2)
    b = 0;
...
if (a != b) {
    ...
}

```

b1 → DSUBUI R3, R1, #2  
b1 → BNEZ R3, L1  
 DADD R1, R0, R0  
L1: DSUBUI R3, R2, #2  
b2 → BNEZ R3, L2  
 DADD R2, R0, R0  
L2: DSUBU R3, R1, R2  
b3 → BEQZ R3, L3

- Prediction accuracy can be improved by accounting for **branch spatial correlations** and looking at the behavior of other branches
  - e.g. in the example above, if the first two branches (b1 and b2) are not taken then the third (b3) will be taken

## A Simple Example: 1-Bit Predictor

```

if (d == 0)
    d = 1;
if (d == 1) {
    ...
}

```

b1 → BNEZ R1, L1  
 DADDIU R1, R0, #1  
L1: DSUBUI R3, R1, #-1  
b2 → BNEZ R3, L2  
L2: ...

Assumption: d alternates between 2 and 0

| d=? | b1 pred. | b1 action | b1 new pred. | b2 pred. | b2 action | b2 new pred. |
|-----|----------|-----------|--------------|----------|-----------|--------------|
| 2   | NT       | T         | T            | NT       | T         | T            |
| 0   | T        | NT        | NT           | T        | NT        | NT           |
| 2   | NT       | T         | T            | NT       | T         | T            |
| 0   | T        | NT        | NT           | T        | NT        | NT           |

A 1-bit predictor that is initialized to not-taken mispredicts all branches

## A Simple Example: Set-Up

```

if (d == 0)
    d = 1;
if (d == 1) {
    ...
}

```

b1 → BNEZ R1, L1  
 DADDIU R1, R0, #1  
L1: DSUBUI R3, R1, #-1  
b2 → BNEZ R3, L2  
L2: ...

| init d    | d==0? | b1        | d before b2 | d==1? | b2        |
|-----------|-------|-----------|-------------|-------|-----------|
| 0         | yes   | not taken | 1           | yes   | not taken |
| 1         | no    | taken     | 1           | yes   | not taken |
| x ≠ {0,1} | no    | taken     | x ≠ {0,1}   | no    | taken     |

**Correlation:** if **b1** is not taken then d is set to 1 and **b2** is also not taken

## A Simple Example: 1-Bit Predictor with 1-Bit Correlation {i.e., a (1,1) Predictor}

```

if (d == 0)
    d = 1;
if (d == 1) {
    ...
}

```

b1 → BNEZ R1, L1  
 DADDIU R1, R0, #1  
L1: DSUBUI R3, R1, #-1  
b2 → BNEZ R3, L2  
L2: ...

Assumption: d alternates between 2 and 0

X/Y: use X if last branch was not taken,  
use Y if last branch was taken

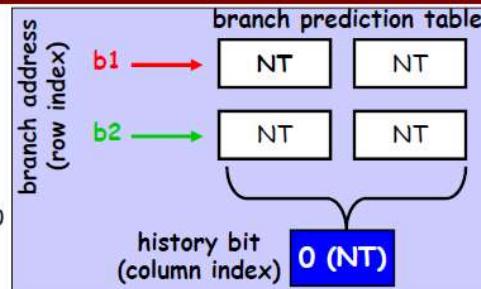
| d=? | b1 pred. | b1 action | b1 new pred. | b2 pred. | b2 action | b2 new pred. |
|-----|----------|-----------|--------------|----------|-----------|--------------|
| 2   | NT/NT    | T         | T/NT         | NT/NT    | T         | NT/T         |
| 0   | T/NT     | NT        | T/NT         | NT/T     | NT        | NT/T         |
| 2   | T/NT     | T         | T/NT         | NT/T     | T         | NT/T         |
| 0   | T/NT     | NT        | T/NT         | NT/T     | NT        | NT/T         |

A (1,1) predictor initialized to NT/NT mispredicts only at the first iter.

## A Simple Example of (1,1) Predictor: Simulation - Cycle 0

```
BNEZ R1, L1
DADDIU R1, R0, #1
L1: DSUBUI R3, R1, #-1
BNEZ R3, L2
...
L2: ...
```

Assumption: d alternates between 2 and 0  
X/Y: use X if last branch was not taken,  
use Y if last branch was taken



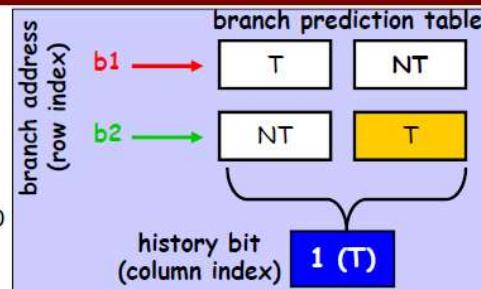
| d=? | b1 pred. | b1 action | b1 new pred. | b2 pred. | b2 action | b2 new pred. |
|-----|----------|-----------|--------------|----------|-----------|--------------|
| 2   | NT/NT    | T         |              | NT/NT    | T         |              |
| 0   |          | NT        |              |          | NT        |              |
| 2   |          | T         |              |          | T         |              |
| 0   |          | NT        |              |          | NT        |              |

A (1,1) predictor initialized to NT/NT mispredicts only at the first iter.

## A Simple Example of (1,1) Predictor: Simulation - Cycle 2

```
BNEZ R1, L1
DADDIU R1, R0, #1
L1: DSUBUI R3, R1, #-1
BNEZ R3, L2
...
L2: ...
```

Assumption: d alternates between 2 and 0  
X/Y: use X if last branch was not taken,  
use Y if last branch was taken



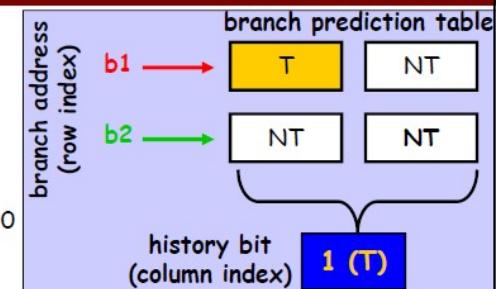
| d=? | b1 pred. | b1 action | b1 new pred. | b2 pred. | b2 action | b2 new pred. |
|-----|----------|-----------|--------------|----------|-----------|--------------|
| 2   | NT/NT    | T         | T/NT         | NT/NT    | T         | NT/T         |
| 0   | T/NT     | NT        |              | NT/T     | NT        |              |
| 2   |          | T         |              |          | T         |              |
| 0   |          | NT        |              |          | NT        |              |

A (1,1) predictor initialized to NT/NT mispredicts only at the first iter.

## A Simple Example of (1,1) Predictor: Simulation - Cycle 1

```
BNEZ R1, L1
DADDIU R1, R0, #1
L1: DSUBUI R3, R1, #-1
BNEZ R3, L2
...
L2: ...
```

Assumption: d alternates between 2 and 0  
X/Y: use X if last branch was not taken,  
use Y if last branch was taken



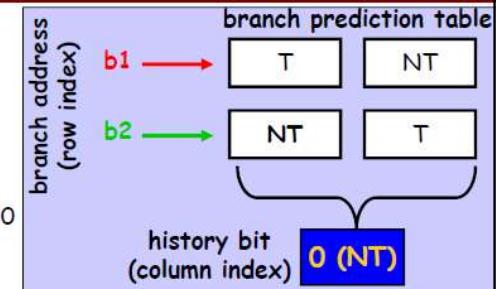
| d=? | b1 pred. | b1 action | b1 new pred. | b2 pred. | b2 action | b2 new pred. |
|-----|----------|-----------|--------------|----------|-----------|--------------|
| 2   | NT/NT    | T         | T/NT         | NT/NT    | T         |              |
| 0   | T/NT     | NT        |              |          | NT        |              |
| 2   |          | T         |              |          | T         |              |
| 0   |          | NT        |              |          | NT        |              |

A (1,1) predictor initialized to NT/NT mispredicts only at the first iter.

## A Simple Example of (1,1) Predictor: Simulation - Cycle 3

```
BNEZ R1, L1
DADDIU R1, R0, #1
L1: DSUBUI R3, R1, #-1
BNEZ R3, L2
...
L2: ...
```

Assumption: d alternates between 2 and 0  
X/Y: use X if last branch was not taken,  
use Y if last branch was taken



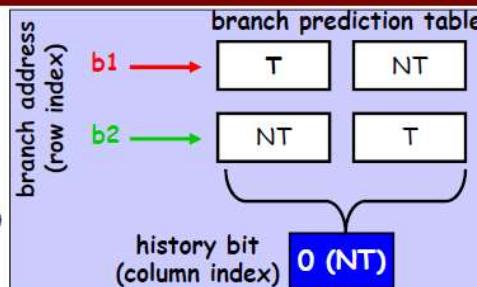
| d=? | b1 pred. | b1 action | b1 new pred. | b2 pred. | b2 action | b2 new pred. |
|-----|----------|-----------|--------------|----------|-----------|--------------|
| 2   | NT/NT    | T         | T/NT         | NT/NT    | T         | NT/T         |
| 0   | T/NT     | NT        |              | T/NT     | NT        |              |
| 2   |          | T         |              |          | T         |              |
| 0   |          | NT        |              |          | NT        |              |

A (1,1) predictor initialized to NT/NT mispredicts only at the first iter.

## A Simple Example of (1,1) Predictor: Simulation - Cycle 4

```
BNEZ R1, L1
DADDIU R1, R0, #1
L1: DSUBUI R3, R1, #-1
BNEZ R3, L2
...
L2: ...
```

Assumption: d alternates between 2 and 0  
X/Y: use X if last branch was not taken,  
use Y if last branch was taken



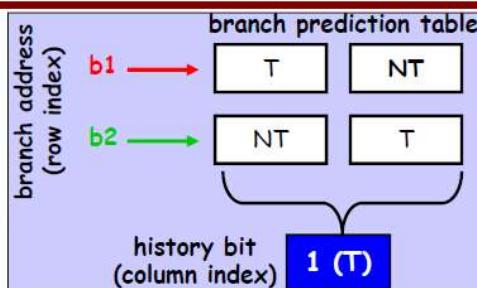
| d=? | b1 pred. | b1 action | b1 new pred. | b2 pred. | b2 action | b2 new pred. |
|-----|----------|-----------|--------------|----------|-----------|--------------|
| 2   | NT/NT    | T         | T/NT         | NT/NT    | T         | NT/T         |
| 0   | T/NT     | NT        | T/NT         | NT/T     | NT        | NT/T         |
| 2   | T/NT     | T         |              | NT/T     | T         |              |
| 0   |          | NT        |              |          | NT        |              |

A (1,1) predictor initialized to NT/NT mispredicts only at the first iter.

## A Simple Example of (1,1) Predictor: Simulation - Cycle 6

```
BNEZ R1, L1
DADDIU R1, R0, #1
L1: DSUBUI R3, R1, #-1
BNEZ R3, L2
...
L2: ...
```

Assumption: d alternates between 2 and 0  
X/Y: use X if last branch was not taken,  
use Y if last branch was taken



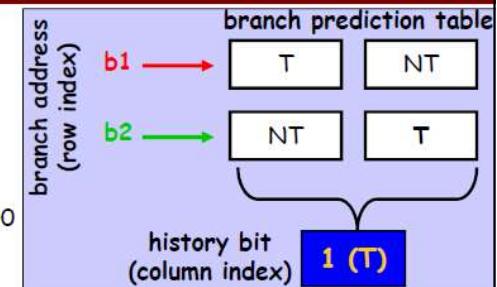
| d=? | b1 pred. | b1 action | b1 new pred. | b2 pred. | b2 action | b2 new pred. |
|-----|----------|-----------|--------------|----------|-----------|--------------|
| 2   | NT/NT    | T         | T/NT         | NT/NT    | T         | NT/T         |
| 0   | T/NT     | NT        | T/NT         | NT/T     | NT        | NT/T         |
| 2   | T/NT     | T         | T/NT         | NT/T     | T         | NT/T         |
| 0   | T/NT     | NT        |              | NT/T     | NT        |              |

A (1,1) predictor initialized to NT/NT mispredicts only at the first iter.

## A Simple Example of (1,1) Predictor: Simulation - Cycle 5

```
BNEZ R1, L1
DADDIU R1, R0, #1
L1: DSUBUI R3, R1, #-1
BNEZ R3, L2
...
L2: ...
```

Assumption: d alternates between 2 and 0  
X/Y: use X if last branch was not taken,  
use Y if last branch was taken



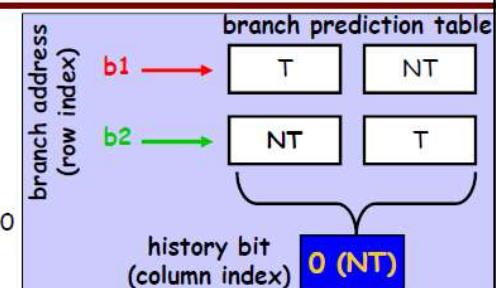
| d=? | b1 pred. | b1 action | b1 new pred. | b2 pred. | b2 action | b2 new pred. |
|-----|----------|-----------|--------------|----------|-----------|--------------|
| 2   | NT/NT    | T         | T/NT         | NT/NT    | T         | NT/T         |
| 0   | T/NT     | NT        | T/NT         | NT/T     | NT        | NT/T         |
| 2   | T/NT     | T         |              | T/NT     | T         |              |
| 0   | T/NT     | NT        |              |          | NT        |              |

A (1,1) predictor initialized to NT/NT mispredicts only at the first iter.

## A Simple Example of (1,1) Predictor: Simulation - Cycle 7

```
BNEZ R1, L1
DADDIU R1, R0, #1
L1: DSUBUI R3, R1, #-1
BNEZ R3, L2
...
L2: ...
```

Assumption: d alternates between 2 and 0  
X/Y: use X if last branch was not taken,  
use Y if last branch was taken



| d=? | b1 pred. | b1 action | b1 new pred. | b2 pred. | b2 action | b2 new pred. |
|-----|----------|-----------|--------------|----------|-----------|--------------|
| 2   | NT/NT    | T         | T/NT         | NT/NT    | T         | NT/T         |
| 0   | T/NT     | NT        | T/NT         | NT/T     | NT        | NT/T         |
| 2   | T/NT     | T         | T/NT         | NT/T     | T         | NT/T         |
| 0   | T/NT     | NT        |              | NT/T     | NT        |              |

A (1,1) predictor initialized to NT/NT mispredicts only at the first iter.

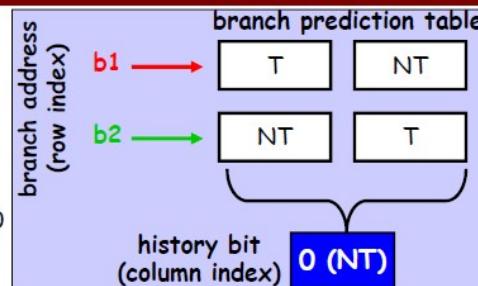
## A Simple Example of (1,1) Predictor: Simulation - Cycle 8

```
BNEZ R1, L1
DADDIU R1, R0, #1
L1: DSUBUI R3, R1, #-1
BNEZ R3, L2
...
L2: ...

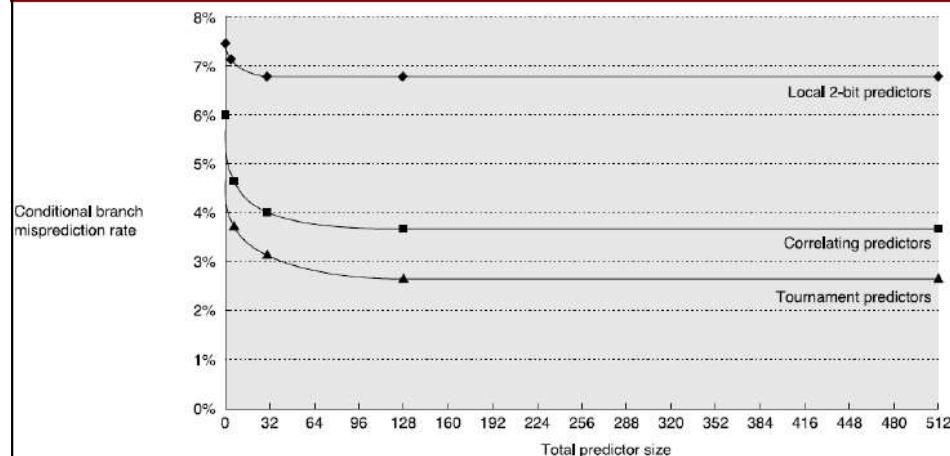
Assumption: d alternates between 2 and 0
X/Y: use X if last branch was not taken,
use Y if last branch was taken
```

| d=? | b1 pred. | b1 action | b1 new pred. | b2 pred. | b2 action | b2 new pred. |
|-----|----------|-----------|--------------|----------|-----------|--------------|
| 2   | NT/NT    | T         | T/NT         | NT/NT    | T         | NT/T         |
| 0   | T/NT     | NT        | T/NT         | NT/T     | NT        | NT/T         |
| 2   | T/NT     | T         | T/NT         | NT/T     | T         | NT/T         |
| 0   | T/NT     | NT        | T/NT         | NT/T     | NT        | NT/T         |

A (1,1) predictor initialized to NT/NT mispredicts only at the first iter.



## Measure: Misprediction Rate for Three Predictors on SPEC89 as Total Number of Bits is Increased



- The prediction capability does not improve beyond a certain size

## Branch Prediction in the Alpha 21264 µP (Tournament Predictors)

### Local Prediction

- each table entry in the first-level local-history table is a 10-bit pattern indicating the direction of the selected branch's last 10 executions.
- The local-branch prediction is a prediction counter bit from a second table indexed by the local-history pattern

### Global Prediction

- the path history is a 12-bit pattern indicating the taken/not taken direction for the last 12 executed branches (in fetch order). The 21264's global-history predictor is a table of saturating counters indexed by the path history

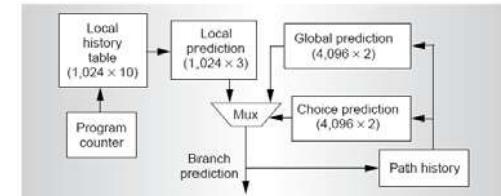


Figure 4. Block diagram of the 21264 tournament branch predictor. The local history prediction path is on the left; the global history prediction path and the chooser (choice prediction) are on the right.

## Branch (Control) Hazards and Their Impact on the Pipeline Performance

- Control hazards can cause a greater performance loss for a pipelined implementation than data hazards. In first approximation:

$$\text{speedupFromPipelining} = \frac{\text{pipelineDepth}}{1 + [\text{branchFrequency} \times \text{branchPenalty}]}$$

- performance loss can vary between 10% and 30% depending on the branch frequency
- generally, the deeper the pipeline, the worse the branch penalty (measured in clock cycles)

- Various strategies to reduce the branch penalty
  - static (compile time) schemes
    - they are fixed for each branch during the entire execution
  - dynamic
    - hardware and software techniques
    - more sophisticated branch prediction schemes

# Exceptions and Interrupts

- “Unexpected” events requiring change in flow of control
  - Different ISAs use the terms differently
- Exception
  - Arises within the CPU
    - e.g., undefined opcode, overflow, syscall, ...
- Interrupt
  - From an external I/O controller
- Dealing with them without sacrificing performance is hard

§4.9 Exceptions

# Handling Exceptions

- In MIPS, exceptions managed by a System Control Coprocessor (CP0)
- Save PC of offending (or interrupted) instruction
  - In MIPS: Exception Program Counter (EPC)
- Save indication of the problem
  - In MIPS: Cause register
  - We'll assume 1-bit
    - 0 for undefined opcode, 1 for overflow
- Jump to handler at 8000 00180



Chapter 4 — The Processor — 41

# Handler Actions

- Read cause, and transfer to relevant handler
- Determine action required
- If restartable
  - Take corrective action
  - use EPC to return to program
- Otherwise
  - Terminate program
  - Report error using EPC, cause, ...



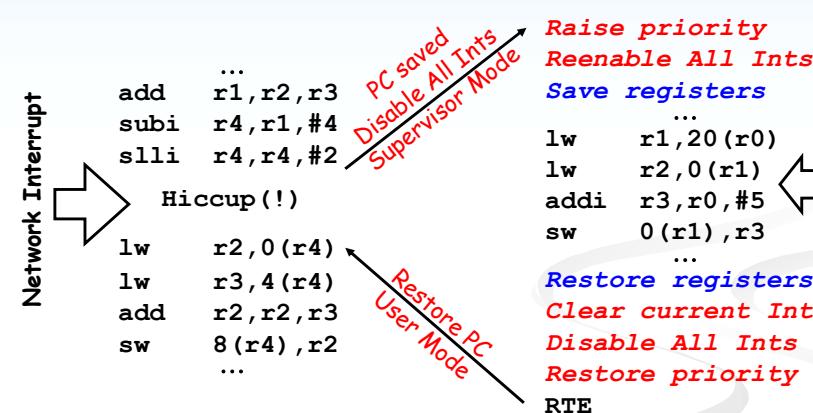
Chapter 4 — The Processor — 43



Chapter 4 — The Processor — 42



# Interrupt Handler



Note that priority must be raised to avoid recursive interrupts!

## Exceptions in a Pipeline

- Another form of control hazard
- Consider overflow on add in EX stage  
add \$1, \$2, \$1
  - Prevent \$1 from being clobbered
  - Complete previous instructions
  - Flush add and subsequent instructions
  - Set Cause and EPC register values
  - Transfer control to handler
- Similar to mispredicted branch
  - Use much of the same hardware



Chapter 4 — The Processor — 45

## Exception Example

- Exception on **add** in
 

```

40    sub  $11, $2, $4
44    and  $12, $2, $5
48    or   $13, $2, $6
4C    add  $1, $2, $1
50    slt   $15, $6, $7
54    lw    $16, 50($7)
...
      
```
- Handler
 

```

80000180  sw    $26, 1000($0) //k0
80000184  sw    $27, 1004($0) //k1
...
      
```



Chapter 4 — The Processor — 47

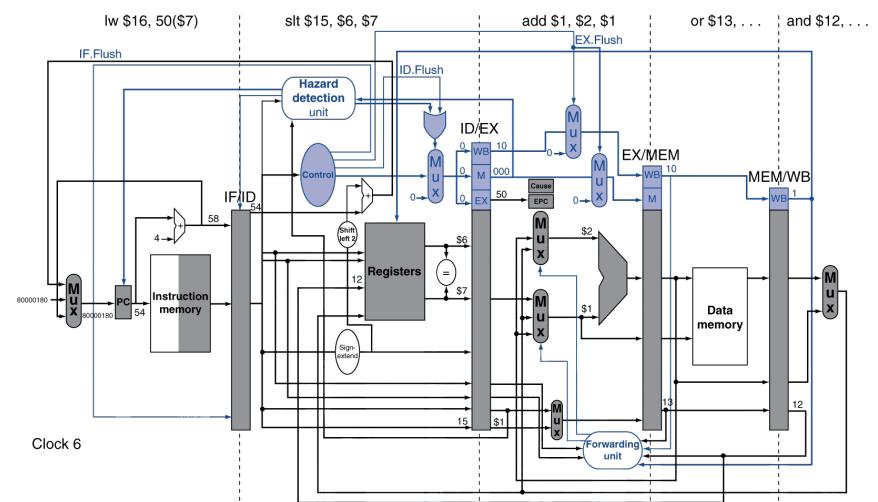
## Exception Properties

- Restartable exceptions
  - Pipeline can flush the instruction
  - Handler executes, then returns to the instruction
    - Refetched and executed from scratch
- PC saved in EPC register
  - Identifies causing instruction
  - Actually PC + 4 is saved
    - Handler must adjust



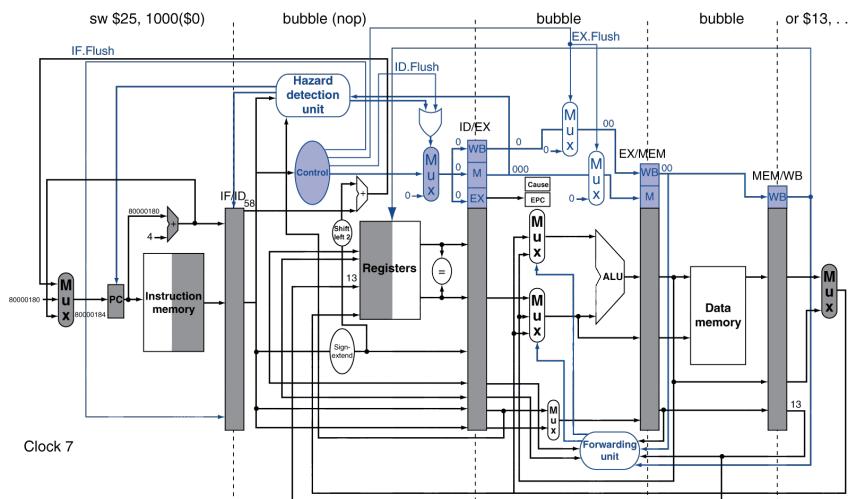
Chapter 4 — The Processor — 46

## Exception Example



Chapter 4 — The Processor — 48

## Exception Example



Chapter 4 — The Processor — 49

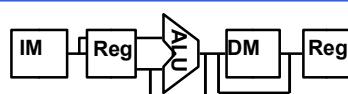
## Multiple Exceptions

- Pipelining overlaps multiple instructions
  - Could have multiple exceptions at once
- Simple approach: deal with exception from earliest instruction
  - Flush subsequent instructions
  - “Precise” exceptions
- In complex pipelines
  - Multiple instructions issued per cycle
  - Out-of-order completion
  - Maintaining precise exceptions is difficult!



Chapter 4 — The Processor — 50

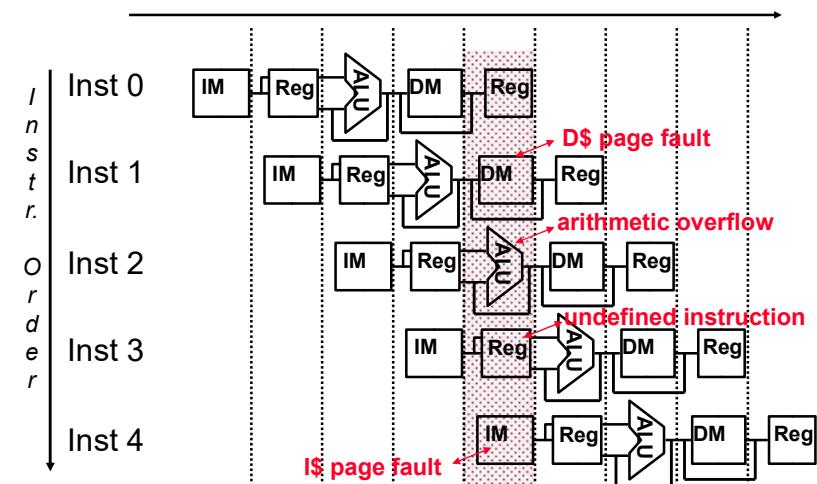
### Where in the Pipeline Exceptions Occur



|                         | Stage(s)? | Synchronous? |
|-------------------------|-----------|--------------|
| ❑ Arithmetic overflow   | EX        | yes          |
| ❑ Undefined instruction | ID        | yes          |
| ❑ TLB or page fault     | IF, MEM   | yes          |
| ❑ I/O service request   | any       | no           |
| ❑ Hardware malfunction  | any       | no           |

❑ Beware that multiple exceptions can occur simultaneously in a *single clock cycle*

### Multiple Simultaneous Exceptions



❑ Hardware sorts the exceptions so that the earliest instruction is the one interrupted first

## Imprecise Exceptions

- Just stop pipeline and save state
  - Including exception cause(s)
- Let the handler work out
  - Which instruction(s) had exceptions
  - Which to complete or flush
    - May require “manual” completion
- Simplifies hardware, but more complex handler software
- Not feasible for complex multiple-issue out-of-order pipelines

